

1 a)

$$f: \mathbb{R}^2 \rightarrow \mathbb{R}, \quad f(x_1, x_2) = x_1^4 + 2x_2^4 - 4x_1x_2$$

$$\begin{cases} f'_{x_1} = 4x_1^3 - 4x_2 & f'_{x_2} = 8x_2^3 - 4x_1 \\ f''_{x_1x_1} = 12x_1^2 & f''_{x_2x_2} = 24x_2^2 & f''_{x_1x_2} = -4 \end{cases}$$

$$\nabla f(\vec{x}) = \begin{bmatrix} 4x_1^3 - 4x_2 \\ 8x_2^3 - 4x_1 \end{bmatrix}$$

$$\nabla^2 f(\vec{x}) = \begin{bmatrix} 12x_1^2 & -4 \\ -4 & 24x_2^2 \end{bmatrix}$$

$$\nabla f(\vec{x}) = 0 \Rightarrow \begin{cases} 4x_1^3 - 4x_2 = 0 \\ 8x_2^3 - 4x_1 = 0 \end{cases} \rightarrow \begin{cases} x_1^3 - x_2 = 0 \\ -x_1 + 2x_2^3 = 0 \end{cases}$$

$$\begin{cases} x_1^3 - x_2 = 0 \\ -x_1 + 2x_2^3 = 0 \end{cases} \rightarrow \begin{cases} x_2 = x_1^3 & (1) \\ -x_1 + 2x_2^3 = 0 & (2) \end{cases}$$

$$(2) -x_1 + 2(x_1^3)^3 = 0 \rightarrow -x_1 + 2x_1^9 = 0$$

$$\rightarrow x_1(2x_1^8 - 1) = 0 \quad \Leftarrow \boxed{x_1 = 0} \Rightarrow \boxed{x_2 = 0}$$

$$2x_1^8 - 1 = 0 \rightarrow x_1^8 = \frac{1}{2} \rightarrow x_1 = \pm \frac{1}{2^{1/8}}$$

$$\boxed{x_1 = \frac{1}{2^{1/8}}} \Rightarrow x_2 = \frac{1}{2^{3/8}}$$

$$\boxed{x_1 = -\frac{1}{2^{1/8}}} \Rightarrow x_2 = -\frac{1}{2^{3/8}}$$

- Gradient is zero when we have the following vectors:

$$\vec{x}_1 = \begin{bmatrix} 0 \\ 0 \end{bmatrix}$$

$$\vec{x}_2 = \begin{bmatrix} 2^{-1/8} \\ -3/8 \\ 2 \end{bmatrix}$$

$$\vec{x}_3 = \begin{bmatrix} -2^{-1/8} \\ -3/8 \\ -2 \end{bmatrix}$$

① matrix $M \in \mathbb{R}^{n \times n}$ is positive definite $\Leftrightarrow x^T M x > 0$
 $\forall x \in \mathbb{R}^n \setminus \{0\}$. This holds if all eigenvalues of M are greater than zero. The matrix is ~~negative~~ negative definite if $x^T M x < 0$, the all eigenvalues must be less than zero.

- Sufficient conditions for local optimum
 If $f: \mathbb{R}^n \rightarrow \mathbb{R}$ is a second order continuous differentiable function:

1. $\nabla f(x^*) = 0$
2. $\nabla^2 f(x^*) > 0 \Rightarrow x^*$ is local minimizer

$$\nabla^2 f\left(\begin{bmatrix} 0 \\ 0 \end{bmatrix}\right) = \begin{bmatrix} 0 & -4 \\ -4 & 0 \end{bmatrix} \rightarrow \text{This matrix is not PD, hence it's not a local minimizer}$$

$$\nabla^2 f\left(\begin{bmatrix} 2^{-1/8} \\ 2^{-3/8} \end{bmatrix}\right) = \begin{bmatrix} \frac{12}{2^{1/4}} & -4 \\ -4 & \frac{24}{2^{6/8}} \end{bmatrix} \rightarrow \begin{matrix} \lambda_1 \approx 16,7 > 0 \\ \lambda_2 \approx 7,7 > 0 \end{matrix}$$

Since both eigenvalues are greater than zero, and the gradient in $x = \begin{bmatrix} 2^{-1/8} \\ 2^{-3/8} \end{bmatrix}$ is zero, we know that the point is a local minimizer.

~~$\nabla^2 f\left(\begin{bmatrix} 2^{-1/8} \\ 2^{-3/8} \end{bmatrix}\right) = \nabla^2 f\left(\begin{bmatrix} 2^{-1/8} \\ 2^{-3/8} \end{bmatrix}\right) \Rightarrow$ The Hessian is the δ~~

$$\nabla^2 f\left(\begin{bmatrix} 2^{-1/8} \\ -2^{-3/8} \\ 2 \end{bmatrix}\right) = \nabla^2 f\left(\begin{bmatrix} 2^{-1/8} \\ -2^{-3/8} \\ 2 \end{bmatrix}\right)$$

\Rightarrow the Hessian is the same for

$$\vec{x}_2 = \begin{bmatrix} 2^{-1/8} \\ -2^{-3/8} \\ 2 \end{bmatrix} \text{ and } \vec{x}_3 = \begin{bmatrix} 2^{-1/8} \\ -2^{-3/8} \\ 2 \end{bmatrix}$$

This means that $\nabla^2 f(\vec{x}_3)$ also has eigenvalues that are greater than zero, which means that it is also a local minimizer.

$\vec{x}_1 = \begin{bmatrix} 0 \\ 0 \\ 0 \end{bmatrix}$ is not a local maximum or minimum since it has one positive and one negative eigenvalue. The Hessian in the point \vec{x}_1 can therefore be neither positive definite nor negative definite.

Answer: $\vec{x}_2 = \begin{bmatrix} 2^{-1/8} \\ -2^{-3/8} \\ 2 \end{bmatrix}$ and $\vec{x}_3 = \begin{bmatrix} 2^{-1/8} \\ -2^{-3/8} \\ 2 \end{bmatrix}$ are local minimizers

$\vec{x}_1 = \begin{bmatrix} 0 \\ 0 \\ 0 \end{bmatrix}$ is neither a local minimizer nor maximizer

$$1b) f: \mathbb{R}^3 \mapsto \mathbb{R}, \quad f(\vec{x}) = \vec{x}^T \cdot A \vec{x} + \vec{b}^T \vec{x}$$

$$A = \begin{bmatrix} -1 & 0 & 1/2 \\ 0 & -1 & 0 \\ 1/2 & 0 & -1 \end{bmatrix} \quad \vec{b} = \begin{bmatrix} 0 \\ 1 \\ 2 \end{bmatrix}$$

$$\frac{\partial f(\vec{x})}{\partial \vec{x}} = (A + A^T) \vec{x} + \vec{b}$$

$$\frac{\partial f(\vec{x})}{\partial \vec{x}} = \begin{bmatrix} -2 & 0 & 1 \\ 0 & -2 & 0 \\ 1 & 0 & -2 \end{bmatrix} \vec{x} + \begin{bmatrix} 0 \\ 1 \\ 2 \end{bmatrix} = \vec{0}$$

$$\rightarrow \begin{cases} \textcircled{1} -2x_1 + x_3 = 0 \\ \textcircled{2} -2x_2 = -1 \\ \textcircled{3} x_1 - 2x_3 = -2 \end{cases} \rightarrow \boxed{x_2 = \frac{1}{2}}$$

$$\textcircled{1} + 2 \cdot \textcircled{3} \Rightarrow -3x_3 = -4 \Rightarrow \boxed{x_3 = \frac{4}{3}}$$

$$x_1 = -2 + 2 \cdot \frac{4}{3} \rightarrow \boxed{x_1 = \frac{2}{3}}$$

$$\text{- Gradient } \nabla f(\vec{x}) = \vec{0} \text{ when } \vec{x} = \begin{bmatrix} 2/3 \\ 1/2 \\ 4/3 \end{bmatrix}$$

$$\nabla^2 f(\vec{x}) = \begin{bmatrix} -2 & 0 & 1 \\ 0 & -2 & 0 \\ 1 & 0 & -2 \end{bmatrix}$$

eigenvalues of $\nabla^2 f(\vec{x})$ are:

$$\lambda_1 = -1, \lambda_2 = -2, \lambda_3 = -3$$

- The hessian is Negative definite in the point $\vec{x} = \left[\frac{2}{3} \ \frac{1}{2} \ \frac{4}{3} \right]^T$ since all it's eigenvalues are less than zero, in the point.
- Because the gradient is also $\nabla f(\vec{x}) = 0$ in the point we know that it's a local maximizer.

Answer: $\vec{x} = \left[\frac{2}{3} \ \frac{1}{2} \ \frac{4}{3} \right]^T$ is a
local maximizer

$$2. \quad f: \mathbb{R}^n \mapsto \mathbb{R}, \quad f(\vec{x}) = \|Ax - b\|_2^2, \quad A \in \mathbb{R}^{m \times n}, \quad b \in \mathbb{R}^m$$

$$\begin{cases} x^{(0)} = a, \quad a \in \mathbb{R}^n & \text{for } t=0 \\ x^{(t+1)} = x^{(t)} - \eta \cdot \nabla f(x^{(t)}), & \text{for } t > 0 \end{cases}$$

$$\frac{\partial f}{\partial x} = 2(Ax - b)^t A$$

$$\begin{cases} x^{(t)} = a \quad \text{where } a \in \mathbb{R}^n & \text{for } t=0 \\ x^{(t+1)} = x^{(t)} - \eta \cdot 2(Ax^{(t)} - b)^t A & \text{for } t > 0 \end{cases}$$

4.

a)

$$\|Aa - b\|_2^2 = (Aa - b)^T (Aa - b)$$

$$= \sum_{i=1}^N (A_{i1}a_1 + A_{i2}a_2 - b_i)^2$$

$$\sum_{i=1}^N (x_i^2 a_1 + y_i^2 a_2 - 1)^2 = \sum_{i=1}^N (A_{i1}a_1 + A_{i2}a_2 - b_i)^2$$

$$\Rightarrow A = \begin{bmatrix} x_1^2 & y_1^2 \\ x_2^2 & y_2^2 \\ \vdots & \vdots \\ x_n^2 & y_n^2 \end{bmatrix} \quad b = \begin{bmatrix} 1 \\ 1 \\ \vdots \\ 1 \end{bmatrix}$$



```
import pandas as pd
import matplotlib.pyplot as plt

csv_file_path = 'HW2_ellipse.csv'

df = pd.read_csv(csv_file_path)
df.columns = ['x', 'y']
x = df['x']
y = df['y']

plt.scatter(x,y)
plt.xlabel('X-value')
plt.ylabel('Y-value')

plt.show()
```




```

import random
from matplotlib import pyplot as plt
import numpy as np
import pandas as pd

def gradient(A:np.array, b:np.array, x:np.array):
    sum1 = 0
    sum2 = 0
    for i in range(999):
        inner_term = (A[i][0] * x[0] + A[i][1] * x[1] - 1)
        sum1 += 2 * inner_term * A[i][0]
        sum2 += 2 * inner_term * A[i][1]
    return np.array([sum1 / 999, sum2 / 999])
    # return 2 * np.dot((np.dot(A, x) - b).T, A)

def learning_rate(A:np.array):
    # 0.012774369968774082
    return abs(0.5 / np.linalg.norm(A))

def get_A(file='HW2_ellipse.csv'):
    df = pd.read_csv(file)
    df.columns = ['x', 'y']
    x = df['x']
    y = df['y']
    x_squared = x ** 2
    y_squared = y ** 2
    # Create a matrix with x_squared as column one and y_squared as column two
    A = np.column_stack((x_squared, y_squared))
    return A

def print_result(A:np.array, b:np.array, x:np.array, dim:int=999):
    print("\nX value : ", x, "\n")
    print("Function-value : ", function(A, b, x), "\n")

def function(A:np.array, b:np.array, x:np.array, dim:int=999):
    sum = 0
    for i in range(dim):
        sum += (A[i][0] * x[0] + A[i][1] * x[1] - 1) ** 2
    return sum

def plot_ellipse(x: np.array, file:str='HW2_ellipse.csv'):
    # Ensure a has two elements
    # Scatter plot from the file
    df = pd.read_csv(file)
    df.columns = ['x', 'y']
    scatter_x = df['x']
    scatter_y = df['y']
    plt.scatter(scatter_x, scatter_y, label='Data Points')

    assert len(x) == 2, "Parameter array must have exactly two elements."

    a1, a2 = x

    # Generate points for the ellipse
    theta = np.linspace(0, 2 * np.pi, 100)
    ellipse_x = np.cos(theta) / np.sqrt(a1)
    ellipse_y = np.sin(theta) / np.sqrt(a2)

    # Plot the ellipse
    # plt.figure()
    plt.plot(ellipse_x, ellipse_y, color='red', label=f'Ellipse: {a1:.2f}x^2 + {a2:.2f}y^2 = 1')

    plt.xlabel('x')
    plt.ylabel('y')
    plt.title('Ellipse and Data Points')
    plt.legend()
    plt.axis('equal')
    plt.grid(True)
    plt.show()

def main():
    dim = 999
    x = np.array([random.uniform(0, 1) for _ in range(2)])
    A = get_A()
    b = np.array([1 for _ in range(dim)])
    stopping_criteria = 1e-5


    while np.linalg.norm(gradient(A, b, x)) > stopping_criteria:
        x -= learning_rate(A) * gradient(A, b, x)

```

```
print_result(A, b, x)
plot_ellipse(x)
```

```
if __name__ == "__main__":
    main()
```


Ellipse and Data Points



4 d)

Yes I think I found the optimal solution for $f(a)$.

$$a = [0.50012765 \ 1.99460575]$$

$$F(a) = 0.46407693961047425$$

Given that $N=999$ and the function value was low I think I found an optimal solution. Adding 999 squared sums of two squared numbers (two numbers that are even multiplied with a_1 and a_2) can easily become a huge number. Our x and y values were in this case relatively small, but the a -terms could easily have made $f(a)$ a huge number if they would have been updated in the ascent direction during gradient descent. Given my small $f(a)$ and the number $N=999$ I think I fit the data well.

Given the convexity of the function we know that the function is monotonically non-decreasing. This means that there could theoretically be another solution that is equally optimal. This task of fitting an ellipse however has a unique solution so I think I found the most optimal solution.

Furthermore the convexity of the function also tells us that the hessian is always greater or equal to zero. This means that if the gradient is zero we know that the function-value won't decrease more.

When I plotted my approximated ellipse next to the scatter plot it is easy to see how similar they are.

I could have got a more accurate result however if I changed my stopping criteria. My stopping criteria was that the norm of the gradient had to be less than " $1e-5$ ", if I would have chosen that as a smaller value I probably would have got a slightly more optimal solution.